

WAKE: a visitation-statistics atlas of the solar neighbourhood — arrival statistics, an exclusion map, and a flyby network on Gaia DR3 kinematics

(Preprint, English version; a Japanese version accompanies this record)

Yukie Maeda
Independent Researcher, Tokyo
ORCID: 0009-0005-3401-9230*

August 2026

Abstract

Should we have been visited? We present the first calculation of visitation statistics for the solar system on the *observed* stellar kinematics of the solar neighbourhood: arrival statistics of close stellar passages on Gaia DR3 (56,286 stars with 6D phase-space data within 100 pc), combined with settlement dynamics in the probe-based, entry-driven convention of Carroll-Nellenback et al. Our contributions are three layers on one catalogue. (1) *Arrival statistics*: a completeness-corrected, exposure-normalized encounter rate with explicitly declared accounting conventions. Reconstructing the accounting of the two published rate camps on our own data, we find their factor-of-two split is dominated by the population that the completeness correction targets, not by data quality; we also identify a third systematic mode (faint-end Gaia RVS radial-velocity artifacts) that evades both value- and error-based quarantines and can inflate encounter rates severalfold, and we introduce a tangential-velocity consistency test to flag it. Our clean rate is $\lambda(1\text{ pc}) = 4.5^{+4.0}_{-2.7}\text{ Myr}^{-1}$ for the FGK+early-M population. (2) *An exclusion map*: conditional statements of the form “if the settled fraction exceeds f^* , the solar system should have been visited within the past 10 Myr (silence rejected at 95%)”; at the standard probe range of 3.07 pc we obtain $f^* = 5.0 \times 10^{-3}$. A three-valued theorem layer imports the survival theory proved in a companion mathematics paper (DOI 10.5281/zenodo.21955413), separating where survival is theorem-guaranteed, where the theorem is silent, and recording that no extinction theorem exists for the full model anywhere. (3) *A flyby network*: on the time-varying graph of close passages, a propellant-minimal visiting route exists at all times in the window, with a single-transfer departure budget of $\Delta v \simeq 5.8\text{ km s}^{-1}$ — a statistical property of the network, not of any particular star. All statements are conditional probability statements over a declared population; undecidable regions are reported as undecidable. This is surveying, not proof.

Contents

1 Introduction

2

*The roles of AI systems and the complete verification protocol are disclosed in Appendix B; this internal verification is not a substitute for human peer review.

2	Data, selection, and quarantine	3
2.1	Catalogue and completeness	3
2.2	Five quarantine systems and a third failure mode	3
3	Propagation and verification	3
4	Arrival statistics and the accounting of encounter rates	4
4.1	Exposure accounting	4
4.2	Three conventions on one data set	5
4.3	Arrival catalogue and structure	5
5	The exclusion map	5
5.1	Definition	5
5.2	The three-valued theorem layer	6
5.3	Front-direction sensitivity	7
6	The flyby network	7
7	Discussion	8
8	Conclusions	8
A	GJ 710, published-value splits, and reconstruction limits	9
A.1	GJ 710 across DR3-era analyses	9
A.2	The reconstructed accounting conventions	9
A.3	Points not determinable from the published texts	10
B	Method record and AI disclosure	11
C	Completeness-map bin width as a hidden regularizer	11
D	Reproducibility and data release	11

1 Introduction

Should we have been visited? If settled systems exist in the Galaxy and expand in the cheapest available channel (slow, small, and old probes launched at close stellar passages), then the question of whether the solar system has been reached is not a matter of speculation but of *measured kinematics*: which stars have passed within probe range, when, and at what rate, given the selection function of the survey that observed them. This paper computes those statistics on Gaia DR3 and overlays settlement dynamics on them.

This is the third branch of a visitation-hypothesis research programme (warp-class channels were treated in ASTROGATION, higher-dimensional channels in TESSERA); the present branch is the sub-light channel, for which the solar neighbourhood’s real kinematics are the decisive input. Our claim to novelty is deliberately narrow: *the overlay of settlement dynamics on real-catalogue encounter statistics*. The first layer alone, encounter statistics, is a well-developed field [2, 1, 10, 8]; we treat it as our verification bench, reproducing published anchors before adding anything (Section 3). The mathematical foundations of the settlement-dynamics layer (survival

and extinction theorems for contact processes on ballistically moving points) were established in a companion mathematics paper [12]; here we import only its assumptions map and guarantees. Throughout, the discipline of the project constitution applies: every output statement is a conditional probability statement over a declared stellar population; regions where the data cannot decide are reported as *undecidable*, never silently filled; and the whole is surveying, not proof.

2 Data, selection, and quarantine

2.1 Catalogue and completeness

We use the Gaia DR3 six-dimensional sample within 100 pc: 56,286 stars with astrometry and radial velocities [9]. Selection is corrected by inverse-probability weighting (IPW) over the DR3 RVS selection function [5], with a floor $s_{\text{floor}} = 0.05$: stars with selection probability below the floor are not up-weighted but reported as an *undecidable population* (697 stars), so that the corrected population is explicitly FGK and early-to-mid M. Late-M and fainter populations are excluded from all rates and flagged as undecidable (a bounded-rate extension is reserved but not used here).

2.2 Five quarantine systems and a third failure mode

Individual-event claims are protected by quarantine flags: astrometric quality, proper-motion signal-to-noise, extreme radial velocities ($|RV| > 550 \text{ km s}^{-1}$, with a whitelist path for confirmed hypervelocity stars), and a radial-velocity error cut ($\sigma_{RV} > 20 \text{ km s}^{-1}$ excluded from individual event judgements while still contributing to rates).

During the accounting cross-checks of Section 4 we identified a *third* systematic mode that evades both the value-based and the error-based quarantines: **faint-end RVS radial-velocity artifacts** — stars at $G \gtrsim 14$ with $|RV| = 150\text{--}900 \text{ km s}^{-1}$ but disk-like tangential motion. Genuine halo stars show halo kinematics in *both* components; a disk-like tangential velocity with an extreme radial velocity is the artifact signature (cf. the DR3 RV quality guidance of [11]). We therefore add quarantine bit 5 ($G > 13.5 \wedge |RV| > 150 \text{ km s}^{-1}$; 1,323 stars; threshold scanned over 120/150/200 km s^{-1} with a stable clean component), operated like the hypervelocity bit: excluded from individual event judgements, rates reported *both ways* (clean and suspect-inclusive), manual review by the tangential-velocity consistency test, and a rescue whitelist (empty after review: every top suspect showed the artifact signature, e.g. $RV = -790 \text{ km s}^{-1}$ with $v_{\text{tan}} = 18 \text{ km s}^{-1}$). Left unquarantined, this population would inflate the sub-parsec encounter rate by a factor of ~ 7 (Section 4); we flag this as a general methodological caution for DR3-based encounter studies.

3 Propagation and verification

Orbits are propagated in an axisymmetric Galactic potential (MWPotential2014-class) by direct numerical integration; an independent epicyclic analytic path serves as the formulation-independence counter-path (gate G2). Verification anchors (gate G3) reproduce published encounters before any new statistics are claimed: the Scholz-star flyby [13, 7], GJ 710 (our value 1.294 Myr, 0.0503 pc, in agreement with [6]; the published-value split among DR3 analyses is documented neutrally in Appendix A), and the encounter-frequency anchor [2]. A dual-potential audit (McMillan17 [14] via a spline tabulation whose interpolation error is itself audited, max 3.7×10^{-5} relative) bounds potential systematics; where the two-potential difference exceeds a




Figure 1: Exposure-normalized arrival rate $\lambda(t)$ within 5 pc for the clean population (bars), with the undecidable shadow of events beyond individual horizons (gray). The envelope is the measured evidence horizon of Gaia DR3: rates are only claimable where the catalogue can support them. Reproduce: `python3 src/wake_p5/paper_figs_main.py`.

star’s measurement uncertainty at far epochs, the star’s *effective horizon* is shortened (“ t_{pot} ”, a third horizon component), so that model uncertainty is treated exactly like measurement uncertainty: as undecidability, not as noise to be averaged.

Each star carries an *individual horizon* t_h : the epoch beyond which its position uncertainty (dominated by the radial-velocity error, $1 \text{ km s}^{-1} \approx 1 \text{ pc Myr}^{-1}$) exceeds half the probe range. Statistics are computed in the *control region* $|t_{\text{ph}}| \leq t_h$; events beyond horizons form the undecidable shadow visible in Figure 1 — the measured “evidence horizon” of the catalogue. Two implementation defects found and fixed by the G2 gate during this work (a parabolic-refinement coefficient and a sharp-minimum breakdown, repaired by exact d^2 -interpolation; all G3 anchors re-passed) are disclosed in Appendix B.

4 Arrival statistics and the accounting of encounter rates

4.1 Exposure accounting

A rate is a count divided by an exposure, and the choice of exposure is a convention that must be declared. We normalize each star’s contribution by its own effective exposure $2 \min(T, t_h)$ within the control region. A naive uniform window ($1/2T$) dilutes the rate as T grows past typical horizons; our own first-look value of 11.7 Myr^{-1} at 1 pc arose from exactly that convention (horizon-ignored, uniform-window) and is reproduced exactly under it — a downward-biased accounting, not a different sky. The exposure-normalized clean rates are

$$\lambda(1 \text{ pc}) = 4.49 [1.78, 8.50] \text{ Myr}^{-1}, \quad \lambda(2 \text{ pc}) = 27.5 \text{ Myr}^{-1}, \quad \lambda(5 \text{ pc}) = 137.3 \text{ Myr}^{-1}$$

(95% star-bootstrap interval at 1 pc; population: FGK+early-M, IPW-corrected; suspect-inclusive values are co-reported in the catalogue). The rate scaling departs from the geometric d^2 law ($6.1 \times$ from 1 to 2 pc, $5.0 \times$ from 2 to 5 pc), reflecting heavy-tailed, few-star-dominated contributions and

the distance structure of the IPW weights; the d^2 -law tests of Section 3 live in the BJ-convention reconstruction, a different accounting.

4.2 Three conventions on one data set

The two published 1-pc rate camps ($19.7 \pm 2.2 \text{ Myr}^{-1}$ [2] and $10.6 \pm 4.5 \text{ Myr}^{-1}$ [8]) use different accounting conventions (window, exposure, completeness target, and population). We reconstructed both conventions from the published texts (the five points for [2] and seven for [8] that could not be determined from the texts are listed neutrally in Appendix A) and applied all three to our single data set:

Convention (on our data)	$\lambda(1 \text{ pc}) [\text{Myr}^{-1}]$	Published
Ours (control region, IPW, per-star exposure)	$4.5^{+4.0}_{-2.7}$	—
BJ+18 reconstruction ($G \leq 12.5$ subset)	23.7 ± 2.6	19.7 ± 2.2
FP26 reconstruction (discrete, reliability cut)	21.3	10.6 ± 4.5

The BJ-convention value is reproduced on clean DR3 data — so residual spurious astrometry is unlikely to be the main driver of the published split. The remaining FP26 gap closes when the faint-end RV artifacts of Section 2.2 are excluded (bright-only discrete counting gives $\simeq 7.6$, consistent with their solar value of 9.1). Based on this reconstruction, the split is dominated by the *population that the completeness correction targets*: [2] corrects to the full mock stellar population, [8] does not correct beyond a window factor. Consistently, a two-ledger check on our own data, correcting by per-star exposure and IPW (denominator ledger) versus by a mock-based $C(t, d)$ completeness map (numerator ledger, Appendix C), yields a measured *population bridge* of $5.3\times$ between the FGK+early-M and full-population rates, compared with the $\sim 9\times$ implied by the mock extrapolation of [2]; the difference is a quantitative basis for critiquing mock-based completeness corrections.

4.3 Arrival catalogue and structure

The arrival-event catalogue (2,197 stars with median perihelion within 5 pc inside $|t| \leq 13 \text{ Myr}$) reproduces the well-known close passages HD 7977 (0.037 pc, -2.76 Myr) and GJ 710 (0.052 pc, $+1.29 \text{ Myr}$), and carries all quarantine, horizon, and both-ways rate metadata. The time-resolved rate $\lambda(t)$ shows no kinematically coherent clumps (all apparent clumps are single-star dominated; pairwise $\Delta UVW \geq 45 \text{ km s}^{-1}$), and the small approaching-vs-receding excess ($+2.1\%$ relative) is the textbook solar-apex dipole [16] (approaching fraction 75.4% toward the apex hemisphere vs. 24.1% away); after subtracting the sample mean-field reflex it reduces to $+0.78\%$ (1.9σ) — consistent with no spatial structure.

5 The exclusion map

5.1 Definition

The map’s conditional statement is defined by a Poisson visit count:

$$N_{\text{vis}} = f \cdot \frac{\langle \hat{p} \rangle}{f} \cdot \lambda(R) \cdot T_{\text{eff}}, \quad N_{\text{vis}} \geq 3 \iff P(\geq 1 \text{ visit}) \geq 95\%, \quad (1)$$




Figure 2: The exclusion map in the (R, f) plane at $v_p = 10 \text{ km s}^{-1}$ (constant mark field). Above f^* (solid), silence is rejected at 95% for the declared population; the dashed line is the population-bridge reference. Gray: undecidable ($R > 5 \text{ pc}$); yellow band: d^2 -extrapolated. At the standard probe range $R = 3.07 \text{ pc}$: $f^* = 5.0 \times 10^{-3}$.

with $T_{\text{eff}} = \max(0, 10 \text{ Myr} - R/v_p)$ the support-limited retrospective window (probe travel time deducted), f the settled fraction, $\langle \hat{p} \rangle / f$ the mark-field factor (unity for the constant field; evaluated on the real arrival kinematics for the front family), and $\lambda(R)$ the clean-population rate interpolated between measured anchors (below 1 pc: d^2 extrapolation, flagged; above 5 pc: undecidable). The clean rate is a *lower bound* on the physical encounter rate (probes do not obey Gaia’s selection function), so the “should-have-been-visited” region is conservatively small; a population-bridge reference layer ($\times 5.3$, model-explicit) and the suspect-inclusive layer are co-reported.

At the standard probe range of [4] ($R = 3.07 \text{ pc}$),

$$f^* = 5.0 \times 10^{-3}$$

— if more than half a percent of systems in the declared population are settled (and launch on entry), the solar system should have been visited within the past 10 Myr. At $R = 1 \text{ pc}$, $f^* = 6.8 \times 10^{-2}$ (Figure 2); boundary uncertainty follows the rate interval ($\times 0.53$ – $\times 2.52$).

5.2 The three-valued theorem layer

The settlement-dynamics layer imports the survival theory of [12] with its assumptions, not beyond them. The assumptions map (mixing number $K_{\text{mix}} = w_1 T_s / R \geq 20$; bounded relative density $\bar{\Lambda} \leq 4$ on a speed shell, which the isotropized DR3 distribution satisfies non-trivially at $\bar{\Lambda} = 3.3$ – 3.6) yields a **three-valued legend**: (i) *C2-guaranteed* (assumptions hold and the numerical threshold is exceeded), (ii) *theorem-silent* (assumptions fail; notably the *entire* standard scan domain of [4], $R = 3.07 \text{ pc}$ with $T_s = 0.1$ – 1 Myr , is theorem-silent, and guarantees open only at $T_s \gtrsim 3 \text{ Myr}$), and (iii) the reminder that *no extinction theorem exists for the full model*




Figure 3: Survival threshold on measured DR3 kinematics (scale-matched to the isotropic calibration): the threshold band shifts upward by $\sim 1.5\times$. The exclusion map’s guarantee layer uses the conservative edge ($m > 2.0$).

anywhere (the companion paper proves the SIR-type bound does not extend), so “settlement impossible” is nowhere theorem-backed. The asymmetry is physical: for short-lived settlements the mathematics is silent; the longer settlements live, the more strongly silence constrains them. Where the theorem is silent the only judgement material is the numerical threshold, which we *re-calibrated on the measured kinematics*: the survival threshold shifts from the isotropic calibration $m \approx 1.0\text{--}1.3$ to $m \approx 1.5\text{--}2.0$ (Figure 3) — real anisotropic kinematics are *harsher* to settlement fronts than the isotropic approximation, which gives the pure-isotropy open problem of [12] an observational demand.

5.3 Front-direction sensitivity

For the front-family mark field, the direction average over two priors (uniform sphere vs. galactocentric bias) shifts the map by only 2.5%; the map is prior-robust. The *per-direction* spread is large (52%): arrivals are apex-concentrated, so a settlement front’s effectiveness depends on its direction relative to the solar apex — a direction dependence we report as a finding connecting to the settlement-direction discussion of [18].

6 The flyby network

On the time-varying graph of clean solar passages (646 nodes within horizons, perihelion states extracted by the numerical propagator), we ask the constitution’s third-layer question: what is the propellant-minimal visiting route? A visitor riding any passing star may transfer, by one ballistic cruise leg, to a star whose perihelion is arbitrarily deep in the solar system. The minimal single-transfer departure budget is

$$\Delta v \simeq 5.8 \text{ km s}^{-1}$$

(5.93 km s^{-1} for the deepest target $d \leq 0.1 \text{ pc}$; 5.85 km s^{-1} for $d \leq 1 \text{ pc}$) — and this is a *statistical property of the network*, not of a particular star: over the error-MC ensemble the route exists in 100% of draws with median $5.7\text{--}5.9 \text{ km s}^{-1}$ (departure charged entirely to propellant, since co-moving riders cannot use flybys to accelerate, so the budget is an upper bound). Direct passages




Figure 4: The flyby network’s node set (clean passages; perihelion distance vs. time) and the minimal single-transfer route into $d \leq 0.1$ pc.

within 0.1 pc in the window are HD 7977 and GJ 710, as in the published record. The graph estimate is validated by a formulation-independent counter-path: Jacobian shooting in the full potential converges to sub-0.1-mpc misses and shifts the transfer budget by at most 0.19 km s^{-1} (tidal terms are subdominant). Multi-leg optimization, full ensemble re-extraction, and real stellar masses are reserved for v1.1 of the network file.

7 Discussion

The philosophy of undecidability. The atlas’s reliability rests on an asymmetry principle (prefer undecidable over falsely decidable), implemented in three mechanisms: individual horizons with three components (measurement growth, potential-model divergence, budget-type truncation), five quarantine systems with both-ways reporting, and the three-valued theorem layer. Nothing in the map is painted by silent extrapolation.

Limitations. The theorem layer judges under an isotropized approximation of the measured velocity distribution (the real distribution is anisotropic; E9 quantifies the direction of the error). The flyby network assumes solar-mass deflectors and single transfers. The d^2 extrapolation below 1 pc and everything above 5 pc are flagged. The population is FGK+early-M; late-M visitation channels are undecidable here (a bounded-rate extension is reserved).

Comparisons with published values rest on our reconstruction of published accounting conventions; the points that could not be determined from the texts are listed openly in Appendix A. DR4 replacement is designed in from the start (the catalogue is a swappable data layer).

8 Conclusions

On the observed kinematics of the solar neighbourhood, with a declared population and declared accounting: (1) the clean arrival rate at 1 pc is $4.5^{+4.0}_{-2.7} \text{ Myr}^{-1}$, and the published factor-of-two split is dominated by completeness-target population choices, not data quality; (2) if more than 0.5%

of systems (standard probe range) are settled, the solar system should have been visited within the past 10 Myr — silence is informative at 95% above that line, and the mathematics guarantees survival of settlement processes only for long-lived settlements, so silence constrains long-lived settlement most strongly; (3) a propellant-minimal visiting route through the real stellar traffic exists at $\Delta v \simeq 5.8 \text{ km s}^{-1}$ at all times. Methodologically we contribute the exposure-accounting analysis, the faint-end RV artifact quarantine with its tangential-velocity test, and the two-ledger completeness check with its measured population bridge. The atlas is surveying, not proof; but it turns the visitation question into numbers that Gaia DR4 can sharpen.

A GJ 710, published-value splits, and reconstruction limits

A.1 GJ 710 across DR3-era analyses

Table 1 collects published DR3-era perihelion parameters for GJ 710 alongside our reproduction. Our numerical integration in the DC95-class potential gives 1.294 Myr and 0.0503 pc; the MWPotential2014-class main path gives 0.0519 pc; and the closed-form linear-motion approximation on the same DR3 inputs gives 1.294 Myr and 0.0513 pc, so the orbit-curvature contribution is at most 0.002 pc in this window.

Analysis	t_{ph} [Myr]	d_{ph} [pc]
This work (DR3 inputs)	1.294	0.0503–0.0519
[6] (N -body)	1.29 (1.26–1.33)	0.052 (0.048–0.056)
[3] (orbit integration)	1.324 ± 0.026	0.052 ± 0.002
[1]	1.292 (1.257–1.334)	0.0636 (0.0595–0.0678)
[8]	1.3446 ± 0.0022	0.0621 ± 0.0023

Table 1: Published GJ 710 perihelion parameters (DR3 era) and our reproduction. Intervals in parentheses are the published 90% ranges.

The perihelion distances fall into two camps: 0.051–0.052 pc ([6], [3], this work) and 0.062–0.068 pc ([1], [8]). The split is not a data-release effect: rerunning our pipeline on DR2 inputs gives 0.0525 pc where [2] published 0.0676 pc, so the same two camps existed in the DR2 generation. The adopted radial velocities differ only in the second decimal (-14.4 vs. -14.42 km s^{-1}), which cannot produce the gap; based on our reconstruction the split is pipeline-specific, and we document it neutrally here rather than adjudicate it.

For the one published perihelion time that stands apart (1.324 ± 0.026 Myr in [3], against 1.29–1.295 Myr in the other analyses and this work, including the distance-consistent linear-method row of the same paper), we record a neutral unit-conversion hypothesis: our 1.294 Myr expressed in units of $\text{pc}/(\text{km s}^{-1})$ equals $1.294/0.97779 = 1.3234$, coinciding with the published value to 0.1%. We state this as an observation about a numerical coincidence, not as an attribution.

A.2 The reconstructed accounting conventions

Table 2 is our reconstruction, from the published texts, of the accounting conventions of [2] and [8], next to ours. It is the basis of the three-convention cross-check of Section 4; every cell is our reading of the texts, and the points we could not determine are listed in Appendix A.3.

	[2]	[8]	This work
Target	the Sun	per-star average over 3,765 systems within 25 pc (the Sun is one sample at the mode of the distribution)	the Sun
Exposure	uniform ± 5 Myr; measured at 5 pc, scaled to 1 pc by $(d/5)^2$	uniform ± 0.47 Myr, derived from the 90th-percentile speed and the 25-pc boundary; direct at $d < 1$ pc	per-star $2 \min(T, t_h)$ with individual horizons; direct at $d < 1$ pc
Time truncation	fractional surrogate counting inside the window; no per-star truncation	uniform $\div 0.70$ window-end correction; no per-star truncation	numerator restricted to the control region $ t \leq t_h$; denominator restricted by the same t_h
Undecidable-type events	spurious solutions filtered; residual bogus encounters retained (acknowledged as an upward effect)	“unreliable” events (1σ exceeding 50% of the nominal distance) excluded; bound systems merged; catalogue spatial completeness (91% at 50 pc) not applied to the rate	$S < s_{\text{floor}}$, out-of-horizon, and window-edge events separated as undecidable
Completeness	two-dimensional $C(t_{\text{ph}}, d_{\text{ph}})$ from a Galaxia mock against a simulated RV selection ($\langle C \rangle \approx 0.14$ at ± 5 Myr); $1/C$ applied to the numerator	uniform 0.70 (window-end only)	IPW $1/S$ over the DR3 RVS selection function with $s_{\text{floor}} = 0.05$
Corrected population	the full stellar population represented by the mock	the RV-bearing catalogue population	the RVS-selectable population (FGK+early-M); the remainder reported undecidable

Table 2: Reconstructed accounting conventions behind the three-convention cross-check.

A.3 Points not determinable from the published texts

Our reconstruction could not determine the following twelve points; they bound the strength of every comparison made in this paper. For [2]: (1) the specific rationale for the ± 5 Myr rate window; (2) whether the Galaxia-based mock includes white dwarfs, brown dwarfs, and unresolved binaries in the completeness normalization; (3) whether losses induced by the astrometric quality filters are intended to be absorbed by $C(t_{\text{ph}}, d_{\text{ph}})$; (4) whether any cuts were applied on radial-velocity values or uncertainties; (5) the bin widths and smoothing form of the completeness map (our reconstruction is sensitive to this choice; Appendix C). For [8]: (1) whether each encounter is attributed to both members of a pair in “per Myr and star”; (2) the denominator of the GJ 710-class rate 0.0210 (our arithmetic gives 0.0236); (3) whether the acceptance and reliability rules apply to the full star–star statistics or only to the solar-encounter analysis; (4) the precise meaning of the ± 4.5 uncertainty; (5) the rationale for not applying the catalogue spatial completeness to

the rate; (6) the relation between the 122 km s^{-1} completeness limit and the 90th-percentile speed used for the window; (7) the reading of the [2] window as “within 15 Myr” in their text, against the $\pm 5 \text{ Myr}$ of the original.

B Method record and AI disclosure

The complete method record is bundled with this record as a repository archive (verification code under `src/`, phase documents and ruling records under `docs/`); this appendix summarizes it and fixes its location.

Verification protocol. Every quantitative claim passed three gates: convergence (G1), formulation independence against an independent analytic counter-path (G2), and invariant anchors that reproduce published encounters before any new statistics are claimed (G3), with a standing rule that anchor failure halts the pipeline. The error ledger is public in the archive. Its two entries most relevant to this paper are: (i) the two propagation-refinement defects found by gate G2 (a vertex coefficient in the parabolic refinement and a sharp-minimum breakdown of the d -parabola, both repaired by exact d^2 -interpolation, after which all G3 anchors re-passed and the full error MC was re-run); and (ii) an episode we record as “agreeing numbers are the ones to distrust”: two mutually confirming rate values (11.7 and 20.3 Myr^{-1}) were both wrong for different reasons (a horizon-ignored exposure convention and the propagation defect above), and the agreement dissolved only under convention identification and engine re-verification.

AI disclosure. The executor of all computations and drafts was an AI system (Claude Code, model Claude Fable 5) operating under the direction and arbitration of the human author, within a written project constitution; the ruling records of every phase gate are in the archive, and the steering-intervention log is empty. This internal verification is not a substitute for human peer review.

C Completeness-map bin width as a hidden regularizer

The mock-based $C(t, d)$ completeness map used in our BJ-convention reconstruction [15, 17] carries a hidden regularization parameter: its bin width. The reconstructed rate scans non-monotonically, $46.5 \rightarrow 23.7 \rightarrow 20.4 \rightarrow 27.1 \rightarrow 20.7 \text{ Myr}^{-1}$ for $\Delta t = 0.5 \rightarrow 3.0 \text{ Myr}$ (Figure 5): fine bins diverge through empty-cell $1/C$ weights, and coarse bins alias. We adopt $(\Delta t, \Delta d) = (1.5 \text{ Myr}, 0.75 \text{ pc})$ with a $\pm 15\%$ grid-systematic band, and co-report the canonical $(1.0, 0.5)$ value. The full scan record is in the bundled archive (`docs/phase2/06-ctd-reference.md`).

D Reproducibility and data release

This appendix is the paper-side anchor of the three-artifact consistency audit among the paper, the released data, and the browser simulator.

Released data. Three JSON files constitute the public data set. `arrival_catalog_v1.json` (schema v1.1) carries the 2,197 arrival events with quarantine, horizon, and both-ways rate metadata, plus per-event astrometry and a 32-surrogate display bundle drawn from the real covariances. `exclusion_map_v1.json` (map version 1.0.1) carries the rate layers, the conditional-statement template in both languages, and the three-valued theorem layer with its numerical-threshold notes. `flyby_network_v1.json` carries the 646 nodes and the minimal-transfer results.




Figure 5: Bin-width scan of the mock-completeness rate reconstruction: a hidden regularization parameter, disclosed.

A manifest (`MANIFEST.json`) records the SHA-256 digest and byte size of each file; the digests of the published versions are frozen in the Zenodo record accompanying this paper.

Regeneration. Every figure, table, and released file regenerates from the archived repository with pinned commands: `python3 src/wake_p2/event_catalog.py` (catalogue), `python3 src/wake_p3/exclusion_map.py` (map), `python3 src/wake_p4/flyby_network.py` (network), `python3 src/wake_p5/paper_figs_main.py` (all five figures), and `tectonic wake_en.tex / tectonic wake_ja.tex` (this document). The catalogue is a swappable data layer; the DR4 replacement path is a rerun of the same commands on the successor catalogue.

Simulator. The browser simulator reads exactly the three released files plus the manifest; hard-coded display values are prohibited by construction, and the schema follows an additive-compatibility principle. Its first artifact is a verification harness of 16 checks that recomputes the paper’s flagship numbers from the released JSONs, including $\lambda(1\text{ pc})$, f^* at the standard probe range, the minimal transfer Δv , the GJ 710 perihelion, and the SHA-256 digest of every data file against the manifest. The harness is the simulator-side anchor of the same audit and is displayed before any visualization. The simulator is served at <https://yukie-lab.github.io/wake-atlas/>, reading the identical released files.

References

- [1] C. A. L. Bailer-Jones. Stars that approach within one parsec of the sun: New and more accurate encounters identified in Gaia data release 3. *The Astrophysical Journal Letters*, 935:L9, 2022.
- [2] C. A. L. Bailer-Jones, J. Rybizki, R. Andrae, and M. Fouesneau. New stellar encounters discovered in the second Gaia data release. *Astronomy & Astrophysics*, 616:A37, 2018.
- [3] V. V. Bobylev and A. T. Bajkova. Search for close stellar encounters with the solar system based on data from the Gaia DR3 catalogue. *Astronomy Letters*, 48:542–549, 2022.
- [4] Jonathan Carroll-Nellenback, Adam Frank, Jason Wright, and Caleb Scharf. The Fermi paradox and the Aurora effect: Exo-civilization settlement, expansion, and steady states. *The Astronomical Journal*, 158:117, 2019.

- [5] Alfred Castro-Ginard, Anthony G. A. Brown, Zuzanna Kostrzewa-Rutkowska, Tristan Cantat-Gaudin, Ronald Drimmel, et al. Estimating the selection function of Gaia DR3 subsamples. *Astronomy & Astrophysics*, 677:A37, 2023.
- [6] Raúl de la Fuente Marcos and Carlos de la Fuente Marcos. An update on the future flyby of Gliese 710 to the solar system using Gaia DR3. *Research Notes of the AAS*, 6:136, 2022.
- [7] Trent J. Dupuy, Michael C. Liu, William M. J. Best, Andrew W. Mann, et al. WISE J072003.20–084651.2B is a massive T dwarf. *The Astronomical Journal*, 158:174, 2019.
- [8] Eloi Fernandez-Puig, Juan Carlos Morales, Ignasi Ribas, Júlia Laguna-Miralles, and Pol Gujosa. Stellar encounters in the solar neighbourhood and the special case of GJ 710. *Astronomy & Astrophysics*, 710:A252, 2026.
- [9] Gaia Collaboration, A. Vallenari, A. G. A. Brown, T. Prusti, et al. Gaia data release 3: Summary of the content and survey properties. *Astronomy & Astrophysics*, 674:A1, 2023.
- [10] J. García-Sánchez, P. R. Weissman, R. A. Preston, D. L. Jones, J.-F. Lestrade, D. W. Latham, R. P. Stefanik, and J. M. Paredes. Stellar encounters with the solar system. *Astronomy & Astrophysics*, 379:634–659, 2001.
- [11] D. Katz, P. Sartoretti, A. Guerrier, P. Panuzzo, G. M. Seabroke, et al. Gaia data release 3: Properties and validation of the radial velocities. *Astronomy & Astrophysics*, 674:A5, 2023.
- [12] Yukie Maeda. Contact processes on ballistic Poisson particles: criticality, fronts, and when motion helps colonization, 2026. Zenodo preprint.
- [13] Eric E. Mamajek, Scott A. Barenfeld, Valentin D. Ivanov, Alexei Y. Kniazev, Petri Väisänen, Yuri Beletsky, and Henri M. J. Boffin. The closest known flyby of a star to the solar system. *The Astrophysical Journal Letters*, 800:L17, 2015.
- [14] Paul J. McMillan. The mass distribution and gravitational potential of the Milky Way. *Monthly Notices of the Royal Astronomical Society*, 465:76–94, 2017.
- [15] Jan Rybizki, Markus Demleitner, Morgan Fouesneau, Coryn Bailer-Jones, Hans-Walter Rix, and René Andrae. A Gaia DR2 mock stellar catalog. *Publications of the Astronomical Society of the Pacific*, 130:074101, 2018.
- [16] Ralph Schönrich, James Binney, and Walter Dehnen. Local kinematics and the local standard of rest. *Monthly Notices of the Royal Astronomical Society*, 403:1829–1833, 2010.
- [17] Sanjib Sharma, Joss Bland-Hawthorn, Kathryn V. Johnston, and James Binney. Galaxia: A code to generate a synthetic survey of the Milky Way. *The Astrophysical Journal*, 730:3, 2011.
- [18] Jason T. Wright, Jonathan Carroll-Nellenback, Adam Frank, and Caleb Scharf. The dynamics of the transition from Kardashev type II to type III galaxies favor technosignature searches in the central regions of galaxies. *Research Notes of the AAS*, 5:141, 2021.